

INDIVIDUAL REFLECTION

Agentic CI System — Final Project | Corporate Intelligence 2025–2026

Carlos Martín Oliveros | 4º Business Analytics & ADE (BA1) | Team 7 | UFV | 25 April 2026

Question 1 — Agent vs. Human Capability (1–5 Scale)

CI Phase	Agent Score	Human Score	Commentary
1. Blind Spot Check	2 / 5	5 / 5	The agent can enumerate common cognitive biases from training data, but genuine
2. Data Collection (OSINT)	5 / 5	3 / 5	The agent's strongest phase. It systematically queries APIs, scrapes structured so
3. AMC Analysis	3 / 5	4 / 5	The agent correctly structures the AMC framework and populates it from retrieved
4a. Scenario Planning	2 / 5	5 / 5	Scenarios require creative synthesis and genuine uncertainty about the future. The
4b. War Game Design	2 / 5	5 / 5	War game design requires anticipating human behaviour, organisational dynamics
4c. Assumption-Based Planning	4 / 5	4 / 5	The agent identified load-bearing assumptions correctly, but the signposts it gener
5. ECOMO / NPV	4 / 5	4 / 5	Arithmetic is reliable; the agent performs NPV calculations and expected value we
6. Early Warning System	4 / 5	3 / 5	Designing KITs and monitoring indicators is well-suited to the agent's systematic m
7. Recommendation	2 / 5	5 / 5	The final recommendation requires integrating all previous phases, weighing contr

The agent adds most value in phases that are systematic and data-intensive (OSINT collection, structured scoring, arithmetic). It is weakest in phases that require genuine strategic creativity, situational judgment, and communicating uncertainty — which are precisely the phases that matter most at the executive decision level. The practical implication is that the agent should be used to accelerate the analytical groundwork so that the human analyst can devote their limited time to the phases where human judgment is truly irreplaceable.

Question 2 — Hallucination Risk

During Phase 3 (AMC Analysis), the agent produced the following statement in its initial output: *'According to Gartner's Magic Quadrant for Cloud AI Developer Services Q1 2026, Google Vertex AI leads in both "Ability to Execute" and "Completeness of Vision", with OpenAI and Anthropic both positioned as Challengers.'* This was presented as a specific citation with high confidence.

■ **Hallucination Detected:** The Gartner Magic Quadrant for Cloud AI Developer Services Q1 2026 was not present in any of the retrieved documents. The agent fabricated both the report name and the specific positioning of the three vendors.

Detection method: The anti-hallucination layer requires all HIGH-confidence citations to have a corresponding retrieved document ID. The Gartner claim had no document ID attached, triggering a mandatory flagging routine. The output was intercepted before being included in the final report and downgraded to MEDIUM confidence with a note: *'Gartner MQ positioning claimed but source document not retrieved — verify independently before citing.'*

Anti-hallucination measures implemented: (1) Mandatory source attribution — every factual claim must reference a document ID from the vector store; claims without a document ID are automatically downgraded to LOW confidence. (2) Confidence scoring — HIGH requires 2+ corroborating sources; MEDIUM requires 1; LOW means the claim is plausible but unverified. (3) Human-in-the-Loop flag — any LOW-confidence finding is presented separately and must be approved by a human analyst before inclusion. These measures caught the hallucination in this case, but they rely on the assumption that fabricated sources are not also fabricated in a way that happens to match a real document ID — a more sophisticated failure mode that remains a risk.

Question 3 — Framework Integrity

The agent applies the CI frameworks correctly at a structural level: it generates AMC tables with the right dimensions, 2x2 scenario matrices with the right axes, and ECOMO calculations with the right formula. The risk of oversimplification is more subtle and operates at the level of *what gets filled in* rather than *how the structure is organised*.

AMC Framework: The Chen & Miller AMC framework requires estimating not just current capability but the *trajectory* of capability — whether a competitor is accelerating or decelerating. The agent rates static snapshots but does not reason about momentum. A human analyst looking at Anthropic's hiring data, funding trajectory, and regulatory positioning would sense that Anthropic's capability rating should be adjusted upward by 0.5–1 point to account for near-term trajectory, a nuance the agent misses because it weights current benchmark performance more heavily than directional signals.

Shell Scenario Method: The Shell method explicitly requires that scenarios be *surprising but plausible* — they should challenge the analyst's existing mental models. The agent's scenarios were internally consistent but not genuinely challenging. Quadrant 4 (Open Source Disruption) was described in terms that would not have surprised anyone following the AI market in 2026. The framework demands scenarios that force the decision-maker to think differently, and that requires the kind of creative provocation that the agent, constrained by its training distribution, cannot reliably produce.

ECOMO: The framework is largely faithfully applied, with the caveat that the agent used generic cost benchmarks (API pricing from vendor websites, industry-average implementation costs) rather than company-specific data. The resulting NPV figures are directionally correct but would require significant adjustment using the focal company's actual cost structure before they could be presented to a CFO as decision-relevant inputs.

Question 4 — AI Tools Used in Development

AI Tool	Used For	Validation Method	Iterations
Claude Sonnet 3.7 (Anthropic API)	Primary agent backbone: AMC scenario, review of generated output	Manual review of generated output	6-8 iterations, mostly for scenario refinement
Claude Claude.ai (this session)	Code generation for Python agent framework	AI-generated code benchmarks	3-4 iterations, mostly for code generation
GitHub Copilot	Autocomplete for boilerplate Python code	Code review and testing	Continuous, mostly for boilerplate code

The most important validation lesson was that AI-generated code is often syntactically correct but semantically wrong — it does what the code says, not what you intended. Every function generated by Claude or Copilot was tested with real data, not just a toy example, before being integrated. Several functions that looked correct in isolation produced wrong results when connected to the RAG pipeline, requiring full rewrites rather than patches.

Question 5 — EU AI Act Risk Classification

Classification: LIMITED RISK (Article 52, EU AI Act 2024/1689) Rationale: The CI agent is not used to make automated decisions that produce legal effects on individuals or significantly affect them. It is an analytical support tool that provides intelligence to human decision-makers who retain full discretion over the final decision. It does not fall into any HIGH-RISK category (Annex III) because it is not used in employment, credit, biometric identification, critical infrastructure, or law enforcement contexts.

Transparency obligations under Article 52: As a limited-risk AI system interacting with human users, the system must inform users that they are interacting with an AI system. In practice, this means every report generated by the agent must include a visible disclosure: *'This analysis was generated with AI assistance. All outputs have been reviewed by a human analyst. Confidence scores reflect source quality at the time of retrieval.'*

What documentation would be required in a real organisation: A technical documentation file describing the system's purpose, design logic, training data (in the RAG sense — the documents indexed), and known limitations. A human oversight protocol specifying which decisions require AI output review before being actable. A logging and audit trail of all agent outputs, enabling post-hoc review if a decision is challenged. A periodic accuracy review — since the agent's performance depends on the quality of its indexed document base, which degrades as the market evolves and documents become stale.

My honest assessment: The most significant gap between our prototype and real-world deployment readiness is not technical — it is governance. The agent produces outputs that look authoritative. Decision-makers under time pressure will be tempted to act on them without the verification steps that the anti-hallucination layer flags as necessary. An organisation deploying this system would need a formal workflow — not just a disclaimer — that enforces human review before agent outputs feed into decisions with material consequences. Without that workflow, LIMITED RISK in legal classification can translate to HIGH RISK in practice.